

RAG Test Document

Use this PDF to test PDF ingestion, chunking, embeddings, FAISS retrieval, and your LangGraph agent.

1. Artificial Project Fact

Project Name: BananaGraph
Created: 2026
Creator: Anshu Kumar
Number of Nodes: 47
Primary Color: Purple
Purpose: BananaGraph is a fictional graph framework created specifically for testing retrieval-augmented generation systems.

2. Artificial Person Fact

Name: Vikas Sharma
Role: Documentation Engineer
Favorite Tool: FAISS
Office: Chandigarh
Special Note: Vikas maintains the fictional BananaGraph documentation.

3. Artificial Product Fact

Product: RAGBox 9000
Version: 7.4.2
Maximum Context: 18,500 tokens
Embedding Model: BananaEmbed-X
Launch Date: 14 August 2026

4. Retrieval Test Facts

Question	Expected retrieved fact
Who created BananaGraph?	Anshu Kumar
How many nodes does BananaGraph have?	47
What color is BananaGraph?	Purple
Who maintains BananaGraph documentation?	Vikas Sharma
What tool does Vikas like?	FAISS
What is RAGBox 9000's version?	7.4.2
What is its maximum context?	18,500 tokens
What is the fictional embedding model?	BananaEmbed-X

5. Negative Test

There is intentionally no information in this document about a person named **Rahul Mehta**. Ask your agent a question about Rahul Mehta and verify that it does not invent an answer from this PDF.

Recommended test: Put this file in your sampledocs folder, run your PDF loader, print the number of pages and chunks, build the FAISS index, then ask: "How many nodes does BananaGraph have?" Your search_docs tool should return the relevant chunk containing **47 nodes**.